Effortless Docker Hosting

How does Hyper.sh compare?

I just want to deploy a Docker application

Manage your own cluster

I just want to deploy a Docker application

Manage your own cluster

K8s/Mesos/Nomad

Let someone else manage your cluster

ECS, GCE, Docker Cloud, Joyent

I just want to deploy a Docker application

Manage your own cluster	Let someone else manage your cluster	There is no cluster
K8s/Mesos/Nomad	ECS, GCE, Docker Cloud, Joyent	Hyper, Triton

How much do you want to do yourself?

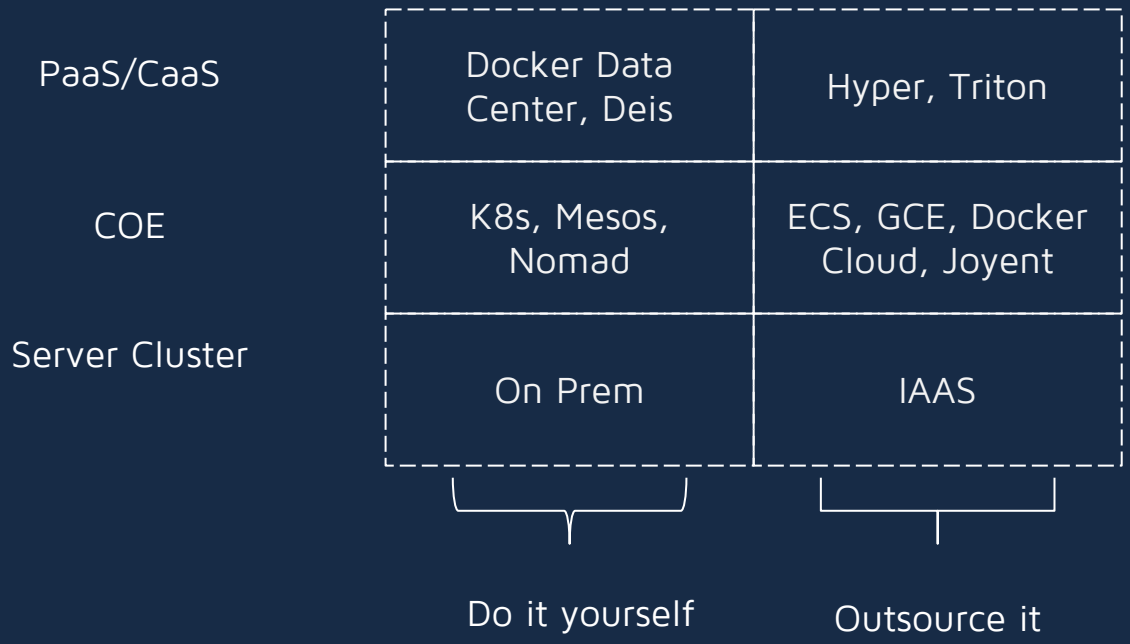

How much do you want to do yourself?

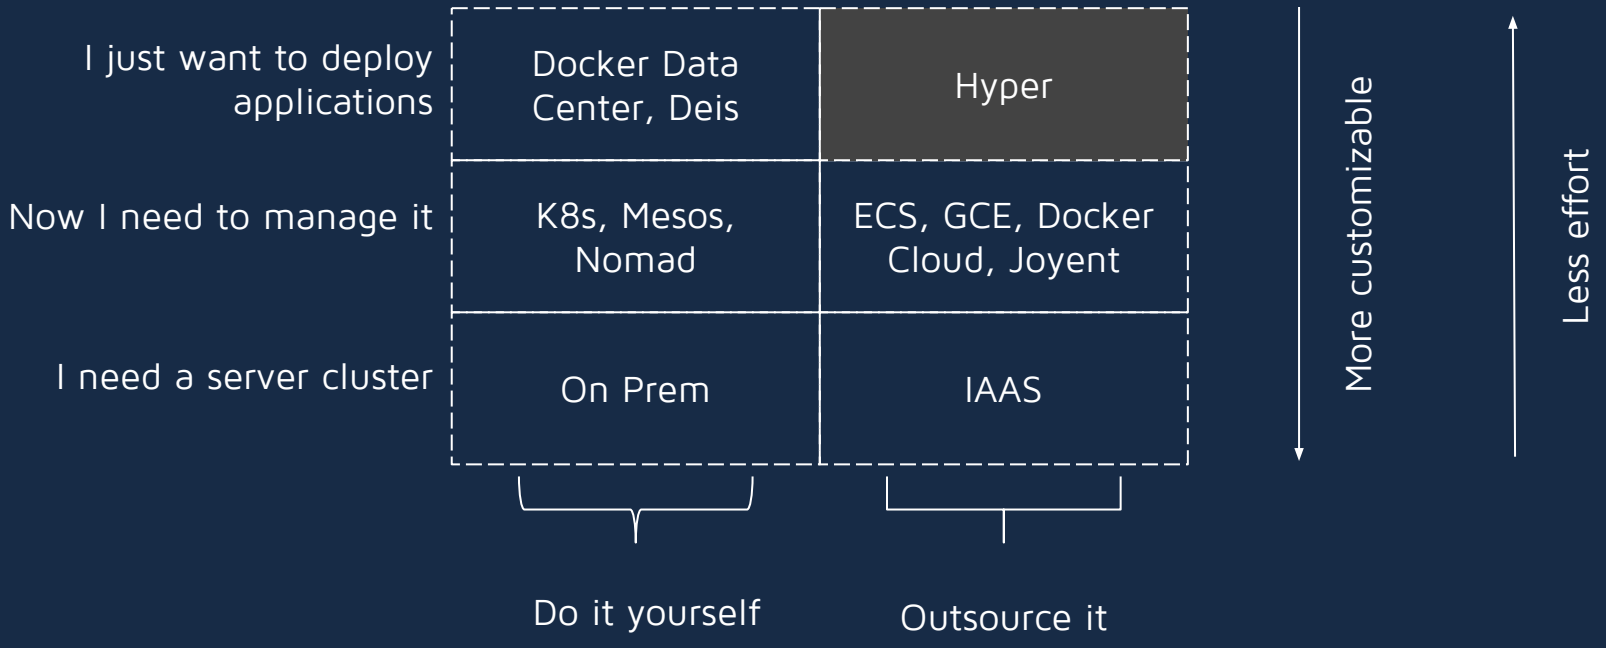

Hyper vs Docker Cloud example

hyper run ubuntu
You're done!

1. Connect to your IAAS provider
2. Provision a VM cluster
3. Now you can deploy an app
4. But, now you still need to maintain your cluster!

DEMO: Deploying a web application with database

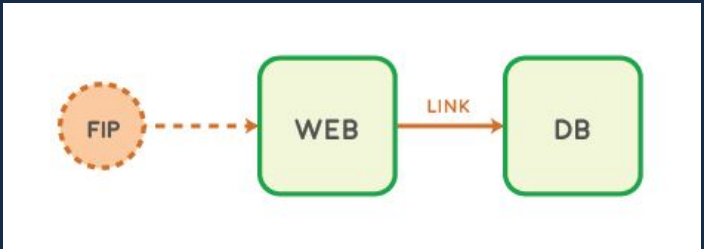


```
hyper run -d --name db hyperhq/postgres
hyper run -d --name web -p 80:80 --link db hyperhq/webapp python app.py
FIP=$(hyper fip allocate 1)
hyper fip attach $FIP web
curl $FIP:80
> Hello: linked database is "tcp://<ip_of_db>"
```

You could also do this in one command with hyper compose as we'll see below

How are people using Hyper.sh?

So how are people using Hyper.sh?

As a host for Docker apps, of course!

hyper compose up -f my-web-app

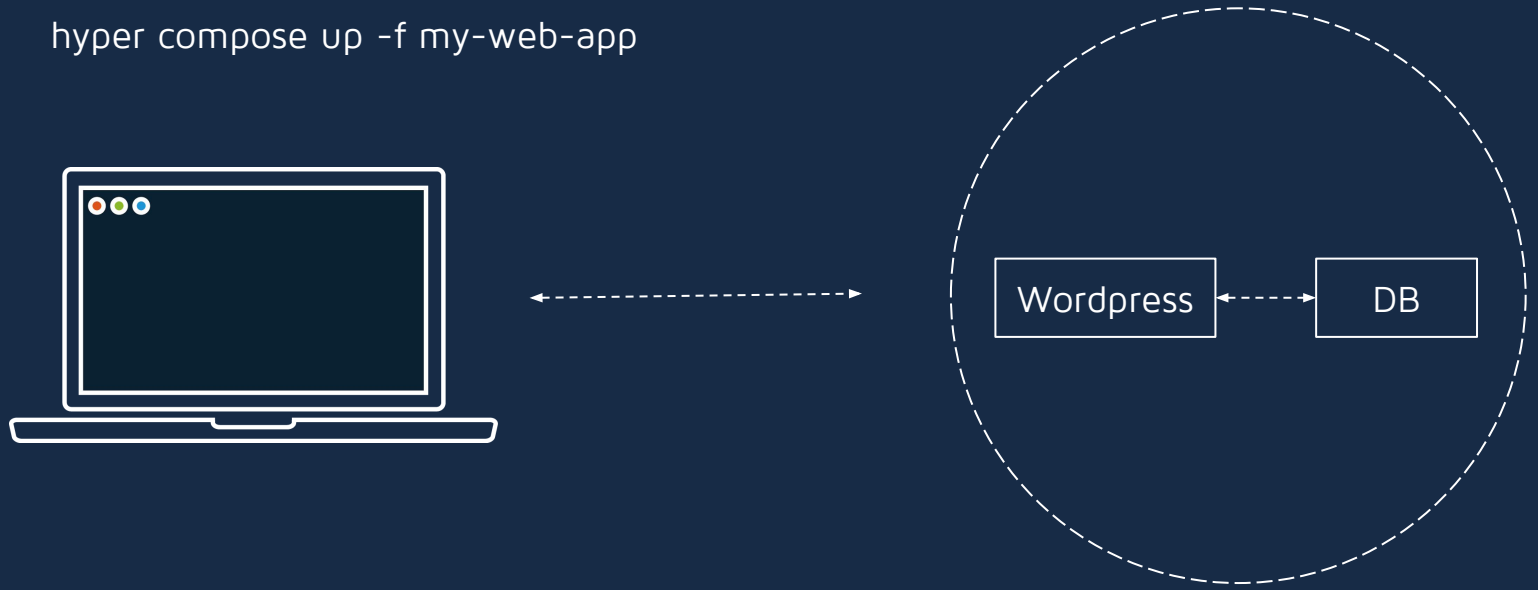

So how are people using Hyper.sh?

Running demos of complex software like Puppet

“Hyper.sh makes it super easy for developers to launch containers in the cloud today.

And the best bit is if you already know the Docker command line tools you already know how to use Hyper”.

Gareth Rushgrove, PuppetLabs/Devops Weekly

So how are people using Hyper.sh?

Development environment with infinite capacity

hyper compose up jupyter-spark.yml

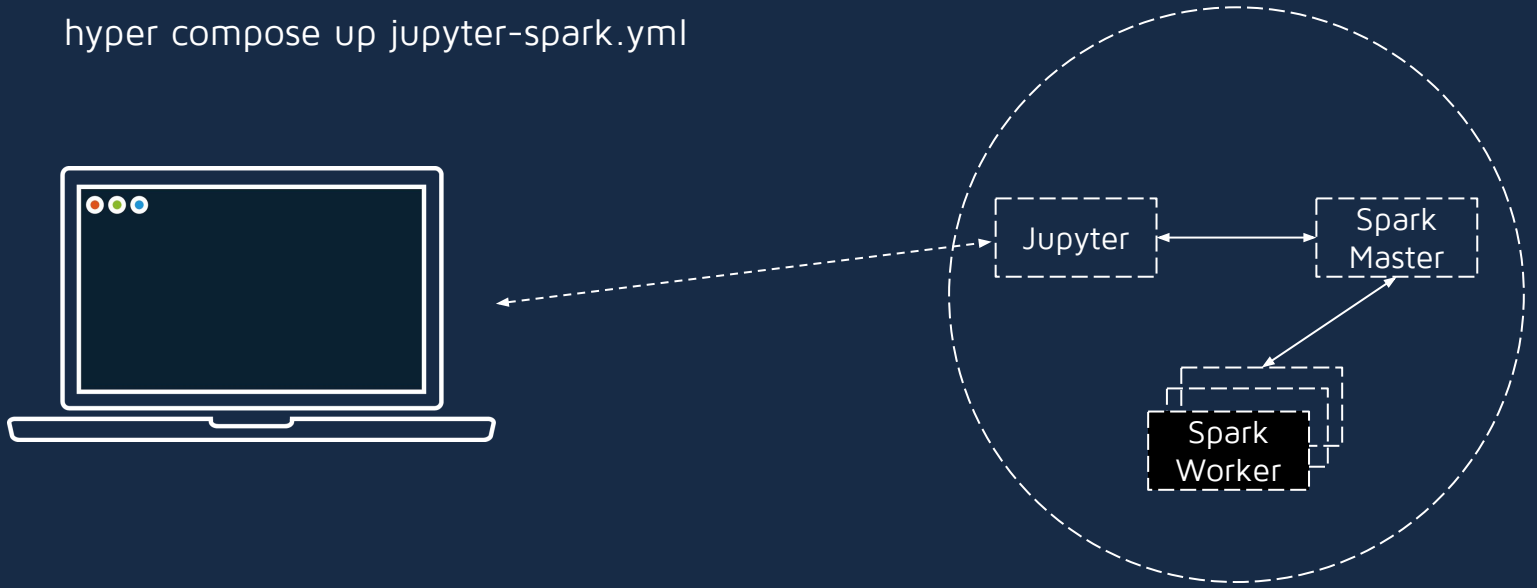

So how are people using Hyper.sh?

'Serverless CI'*

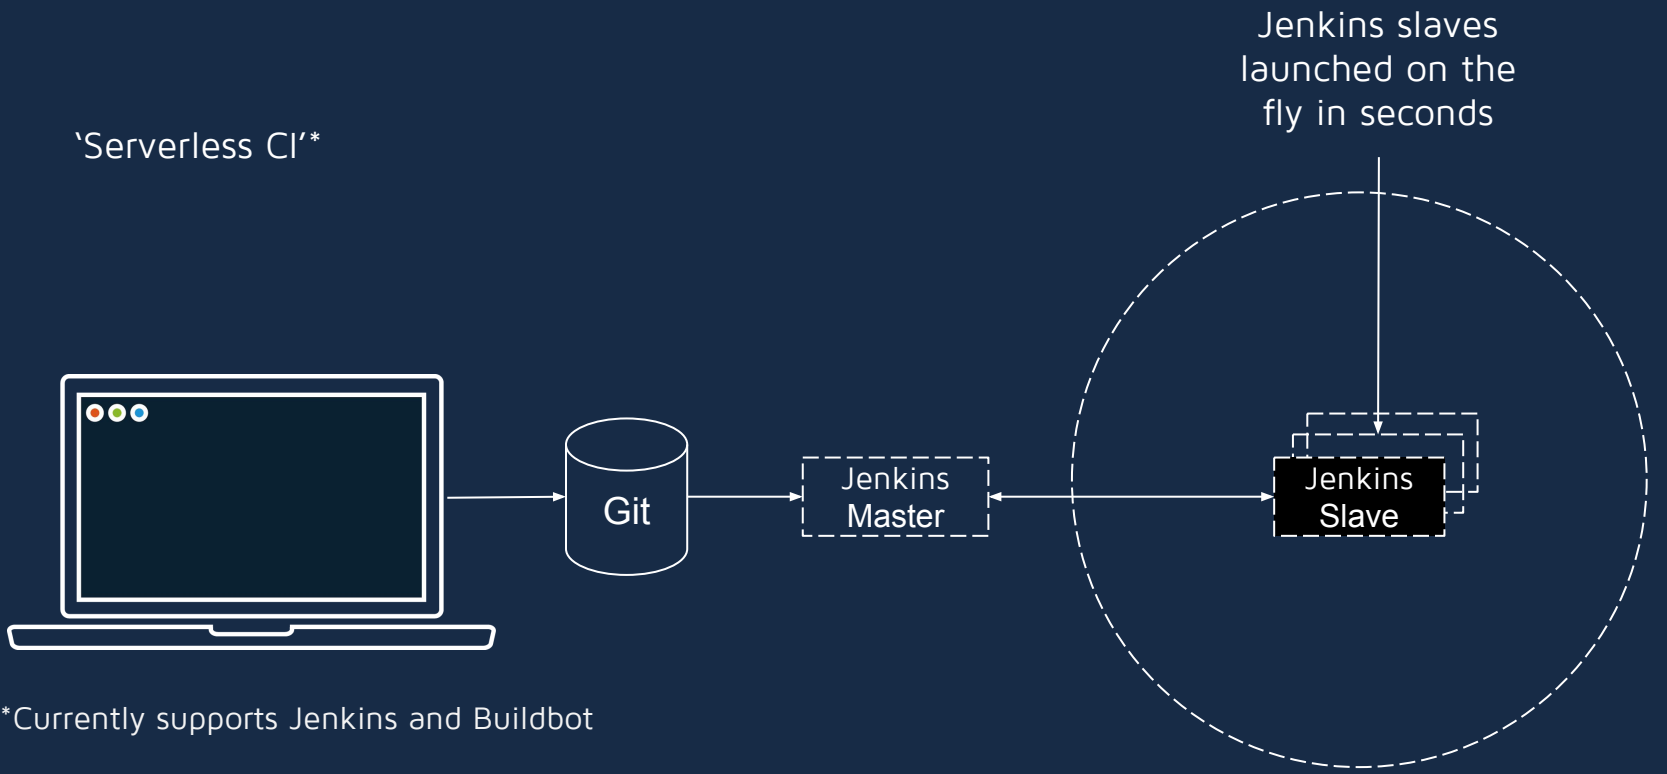

*Currently supports Jenkins and Buildbot

Delivery pipeline

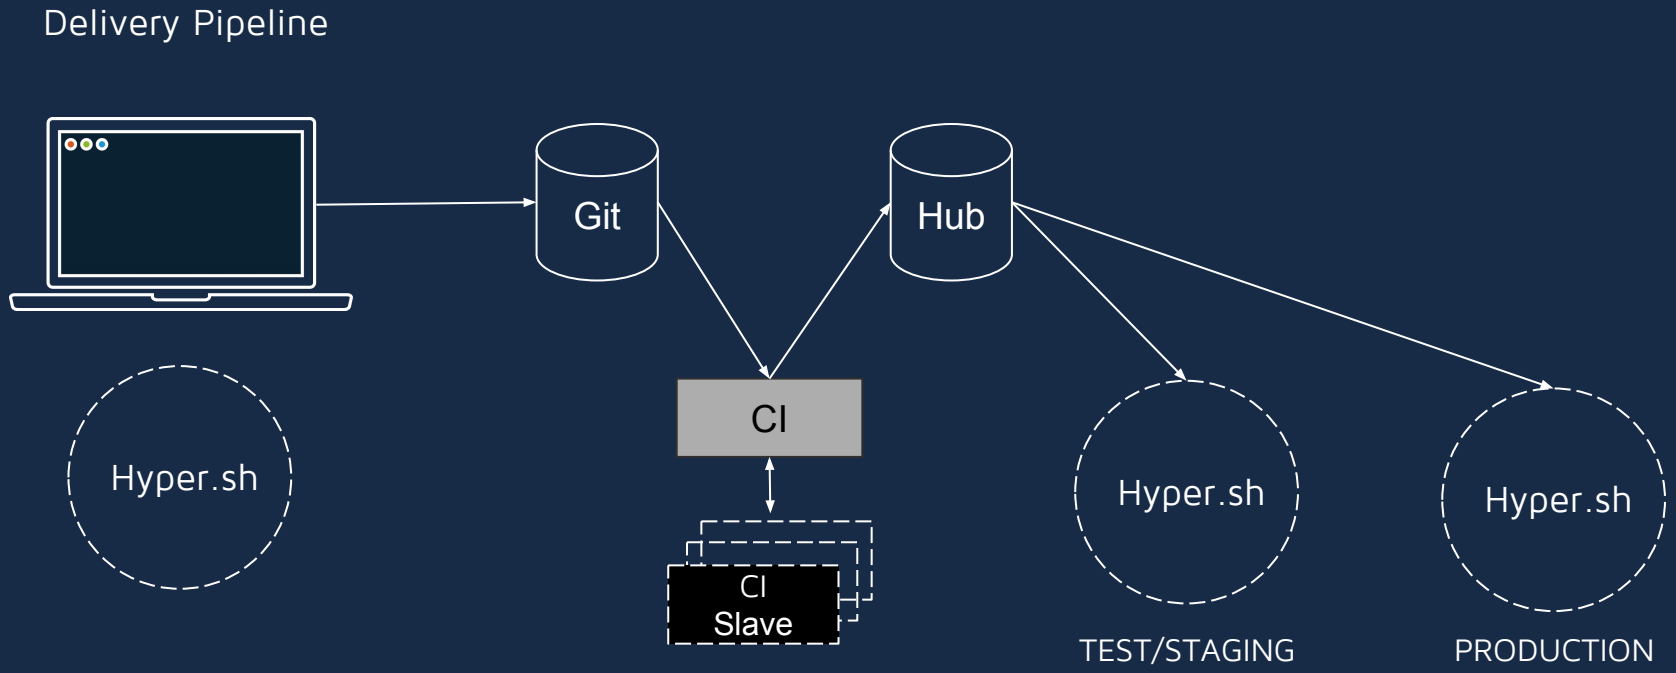

Under the hood

Networking

All your containers can access each other on L2

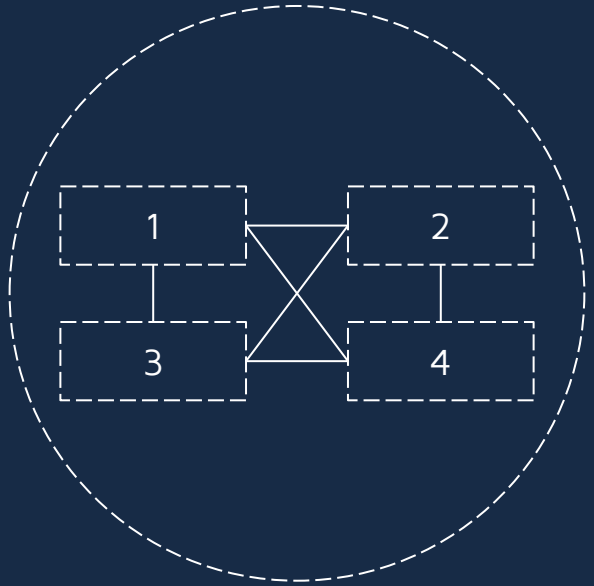

Network testing with iperf yields ~2.7Gbits/sec

Networking

All your containers can access each other on L2

but you can also use security groups to create layers

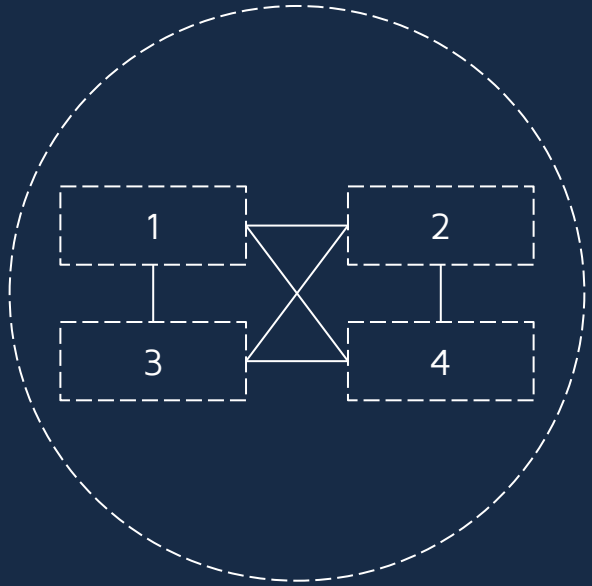

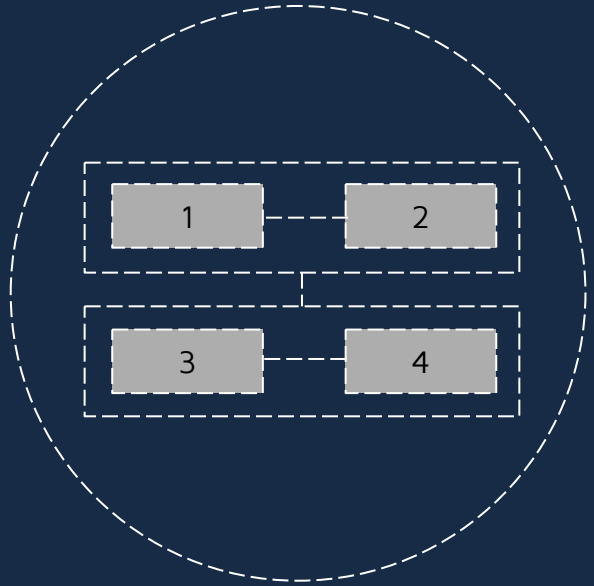

Network testing with iperf yields ~2.7Gbits/sec

Storage

EBS like volume store with automatic replication and backups

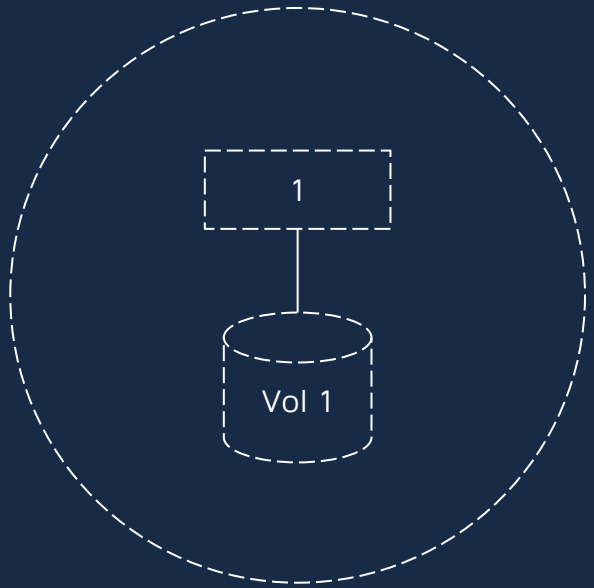

All storage is SSD yielding 250-300MB/s writes

Storage

EBS like volume store with automatic replication and backups
plus snapshots for redeploys

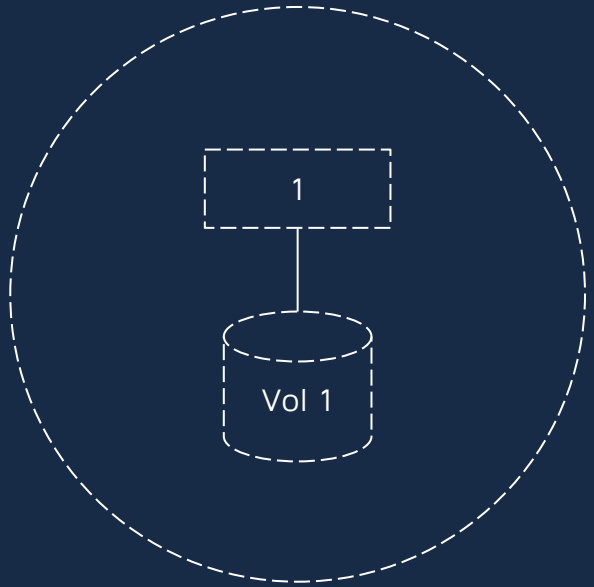

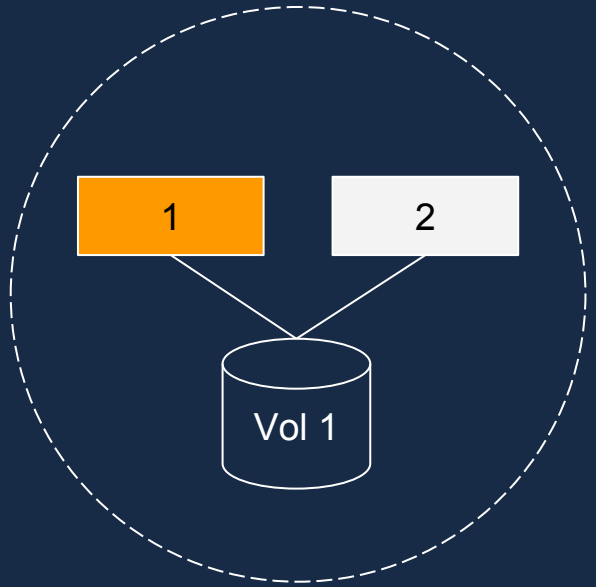

All storage is SSD yielding 250-300MB/s writes

Choose your registry

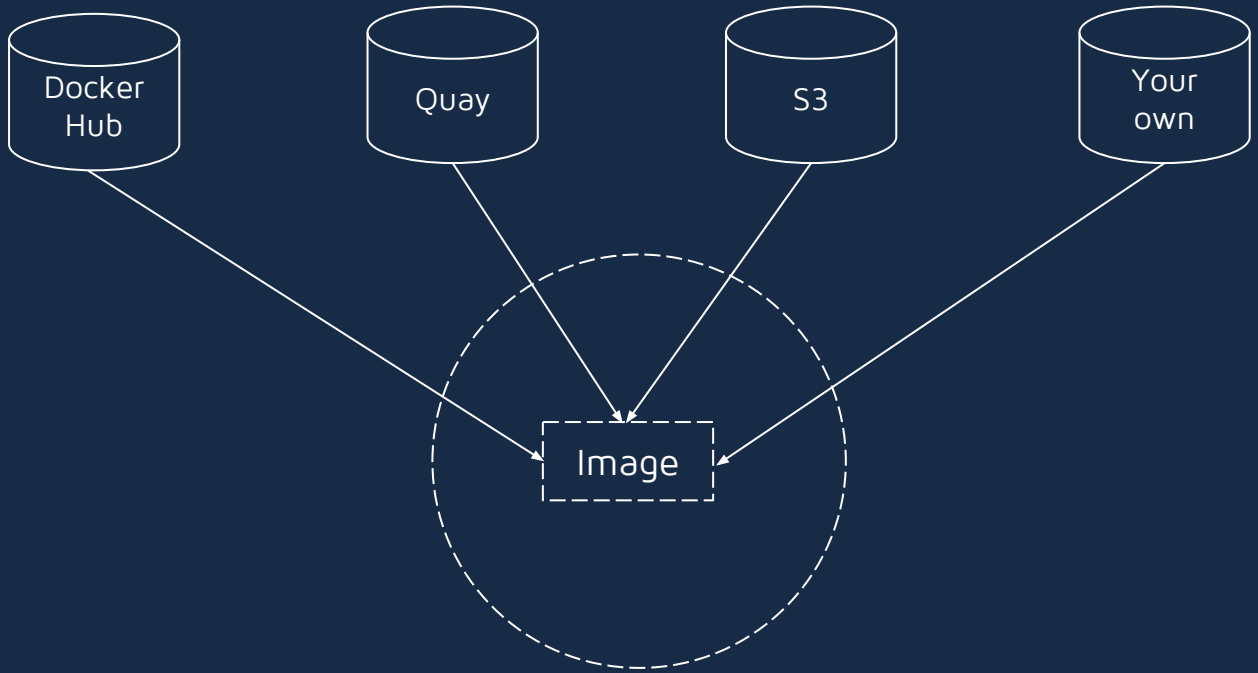

Further information

Further information

Public Roadmap: https://trello.com/b/7fEwaPRd/roadmap

Twitter: https://twitter.com/hyper_sh

Slack: https://slack.hyper.sh/

Blog: https://blog.hyper.sh/

Website: https://hyper.sh/